

Live Demo — Master Document

From Ground Observation to Coordinated Response

For the team — use this for demo-day script writing, the pitch deck, the live walkthrough, and judge Q&A. It describes the automated **/demo** experience exactly as built and verified.

What it is	One scripted incident, driving the Pilgrim, Volunteer and Management screens together on one page — a self-running ~4-minute walkthrough of the whole closed loop.
Where	Route: /demo (the three real apps — /pilgrim, /field, /management — are unchanged)
Runtime	~3 min 40 s end to end, 20 scripted beats. Fully deterministic — no live model decides whether it proceeds.
Controls	Run · Pause · Resume · Skip step · Restart · Exit demo, with a DEMO nn / 20 step counter and progress bar.
Data	Every value is synthetic. Labelled “Data: Synthetic / Simulation”; the pilgrim panel is labelled “demo scenario — not a real pilgrim.”
Branch	feat/live-demo-mode (pushed). Not merged to main. main stays at the known-good build.

1 · The story in one breath

A Tamil-speaking pilgrim reports that her child is missing. The Kumbh Setu Assistant guides her in her own language. A volunteer nearby uses Setu AI to translate the situation and file a structured lost-child ground report — after a human confirmation. A second volunteer, independently, files a matching report. The control room sees two independent reports, corroborates them, and Kumbh Pulse raises one explainable emerging signal for Ghat 4. The operator opens the “why is this zone at risk?” breakdown, promotes the signal to a Critical incident, and asks Ops Copilot for a responder. Ops Copilot recommends V-233 — nearest, crowd-trained, with a Tamil translation bridge. The operator confirms the dispatch. V-233 accepts, arrives, and resolves — the child is reunited. The Event Log shows the entire chain, with an actor and a timestamp on every step.

One incident. Three roles. One connected response loop.

The response loop

Observe	A pilgrim or volunteer reports something on the ground.
Corroborate	Independent reports are cross-checked, not taken at face value.
Signal	Kumbh Pulse raises one explainable signal worth a human's attention.
Decide	The control room promotes it to an incident and chooses a responder.
Resolve	A volunteer accepts, arrives and closes the loop on the ground.
Audit	Every step is logged with actor and time in the Event Log.

Two lines to land at the end

- “Setu AI helps people understand and act.”
- “Kumbh Pulse helps the control room understand what is happening.”
- Close on: **KUMBH SETU — connected response infrastructure for mass gatherings.**

2 · Full script — 20 beats

Columns: beat, on-screen clock, active panel, what the audience sees, and the voiceover line to say over it. The clock is the story timeline shown in the control bar, not real elapsed time.

#	Clock	Panel	On screen	Voiceover
1	0:00	Intro	Title card: three role cards (Pilgrim / Volunteer / Management) and the Observe->Corroborate->Signal->Decide->Resolve->Audit loop.	"Kumbh Setu is one connected response platform for mass gatherings. One incident — watch it move across all three roles at once."
2	0:20	Pilgrim	Pilgrim panel, labelled 'demo scenario — not a real pilgrim'. Tamil message "என் குழந்தையை காணவில்லை" with the gloss "My child is missing."	"A pilgrim reports, in Tamil, that her child is missing. She does not have to switch languages to ask for help."
3	0:35	Pilgrim	Kumbh Setu Assistant replies (emergency styling): stay where you are, nearest Help Desk is Ghat 4 Control, asks age and clothing, says a volunteer is being alerted. Provenance line: 'Powered by Sarvam · Grounded in Kumbh Setu data'.	"The assistant answers in her language, gives lost-child guidance, and stays with her while help is arranged."
4	0:50	Volunteer	Volunteer panel activates. "Help me talk to a Tamil-speaking pilgrim" -> Setu translates: "My 6-year-old son is missing. He was near the Ramkund steps about ten minutes ago. He is wearing a red shirt." (Tamil->English · phrasebook · 96%).	"A volunteer nearby bridges the language gap. Setu translates the situation both ways."
5	1:05	Volunteer	Volunteer: "Log this as a lost child report: 6-year-old boy in a red shirt, missing near Ramkund." Setu AI shows a structured report (category, zone, description, severity High, source). Amber card: 'Setu proposed: File ground report — high-impact action · human confirmation required'.	"Setu structures the report — but it does not file anything on its own. A human has to confirm."
6	1:20	Volunteer	The proposal flips to 'Confirmed by a human — action executed'. Green card: GR-4001 submitted to the control room.	"The volunteer confirms. Report GR-4001 goes to the control room."
7	1:35	Volunteer	Second volunteer (R. Kamble · V-241): "I also have a report — a lost boy about 6 near Ramkund. Someone found a child crying by Gate 3." Another 'Setu proposed: File ground report'.	"Independently, a second volunteer has the same picture from a different spot."
8	1:50	Volunteer	Confirmed. Green card: GR-4002 submitted — control room now has two.	"Two independent reports, filed separately, now sit in the control room."
9	2:05	Management	Management panel activates. Field reports list: GR-4001 and GR-4002, same area, same description.	"The control room isn't reacting to one shout. It has two independent accounts."
10	2:20	Management	GR-4001 marked 'corroborated'. Chain: 2 independent reports -> Corroboration -> Emerging signal. Amber card: 'Multiple lost-person reports — Ghat 4', confidence 69%, 'Kumbh Pulse decision support · synthetic'.	"The two reports corroborate each other. Kumbh Pulse raises one emerging signal — not fifty alerts, one."
11	2:35	Management	Kumbh Pulse panel: Ghat 4, score 84/100, band RED, forecast Critical in 10–20 min, narrative about a missing child on the Ramkund approach during pre-aarti density.	"Ghat 4 moves into the red band. This is decision support — it never acts by itself."
12	2:45	Management	'Why is this zone at risk?' breakdown opens: Crowd Activity 42%, Movement Change 26%, Checkpoint Reports 18%, Historical Pattern 14%; supporting signals listed; active incident noted; 'Data: Synthetic / Simulation'.	"And it shows its working: the four factors behind the score, and the signals behind those. Explainable, not a black-box number."

#	Clock	Panel	On screen	Voiceover
13	3:00	Management	Operator promotes the corroborated signal. Incident created: KS-1003 · Lost Person · Critical · Ghat 4 (triaged).	"The operator promotes it to a Critical incident — one deliberate action, fully logged."
14	3:20	Management	Ops Copilot: prompt 'Recommend a responder — the family only speaks Tamil.' Reasoning: V-233 available, ~120 m, crowd-trained, shortest ETA; Tamil bridge needed; hold V-218 for medical. Recommendation, then 'Setu proposed: Assign volunteer — human confirmation required'.	"Ops Copilot reasons over the operational picture and recommends V-233, with a translation bridge for the family. It proposes — it doesn't dispatch."
15	3:45	Management	Operator confirms. Green: V-233 dispatched — task created, volunteer notified. Incident card shows Responder V-233 · ETA ~4 min.	"The operator confirms the dispatch."
16	4:00	Volunteer	Volunteer panel switches to the task view: Ghat 4 · KS-1003 · CRITICAL -> Accepted — en route. Progress checklist: Dispatched done, Accepted done.	"V-233 gets a critical task and accepts — en route."
17	4:15	Volunteer	Arrived on site — checklist advances.	"Arrives on site at Ghat 4."
18	4:30	Volunteer	Resolved — 'Child reunited — loop closed'. All four checklist states done. Management incident shows Resolved.	"Resolved. The child is reunited with the family, and the incident closes."
19	4:40	Management	Event Log — audited chain, ten rows, each with actor and time: report created ×2, corroborated, promoted, incident reported, triaged, dispatched, accepted, arrived, resolved.	"And here is the whole chain — every step, who did it, and when. That is the audit trail a real operation needs."
20	4:55	Outro	Closing card: the full loop, then 'One incident. Three roles. One connected response loop.' / 'Setu AI helps people understand and act.' / 'Kumbh Pulse helps the control room understand what is happening.' / KUMBH SETU.	"One incident, three roles, one connected loop. Setu AI helps people understand and act; Kumbh Pulse helps the control room understand what's happening. Kumbh Setu — connected response infrastructure for mass gatherings."

3 · Pitch-deck outline (maps to the demo)

A 10-slide spine. Slides 4–8 are the live demo itself; if the demo runs, those slides are just holding frames.

1	Title	Kumbh Setu — connected response infrastructure for mass gatherings. One line, one image.
2	The problem	At Kumbh scale, reports come from everywhere in many languages; the gap is not data, it's a connected, corroborated, auditable response. Name the failure mode: crowd crush from uncoordinated response.
3	The idea	Three roles, one loop: Observe -> Corroborate -> Signal -> Decide -> Resolve -> Audit. Setu AI for people; Kumbh Pulse for the control room.
4	Demo — Pilgrim	Tamil lost-child report + assistant guidance in-language. (Demo beats 2–3.)
5	Demo — Volunteer	Translate -> structured ground report -> human confirm -> GR filed. A second independent report. (Beats 4–8.)
6	Demo — Corroboration & Pulse	Two reports -> one explainable emerging signal -> Ghat 4 red -> 'why is this zone at risk?' factor breakdown. (Beats 9–12.)
7	Demo — Decide & Dispatch	Promote to Critical incident -> Ops Copilot recommends a responder -> human confirm -> dispatch. (Beats 13–15.)
8	Demo — Resolve & Audit	Accept -> arrive -> resolve; Event Log shows the full chain with actor + time. (Beats 16–19.)
9	How it's built / what's real	Sarvam is genuinely integrated (chat 105B, Mayura translate, Bulbul TTS, native tool-calling) in the real apps. The demo is deterministic and offline-safe. No database; human confirmation gates on every high-impact action.
10	Ask / close	One incident, three roles, one loop. What we need next (pilot zone, data partners). Close on the wordmark.

4 · Demo-day runbook

Before you present

- Open **/demo** in a full browser window on the projector display. Nothing else needs to be open.
- Do a full dry run once (Run -> let it finish). Then leave it on the pre-start title card.
- Have the three real apps ready in separate tabs only as a backup: `/pilgrim`, `/field?tab=tasks`, `/management`.
- Confirm network: Sarvam status at **/api/setu/status** should return all true. The demo still runs fully if it doesn't — it never calls Sarvam.

Driving it

- **Run Live Demo** — starts the 20-beat timeline; it advances on its own (~3 min 40 s).
- **Pause / Resume** — freezes on the current beat; everything on screen stays put. Use it to talk over a beat.
- **Skip step** — jumps to the next beat immediately (useful if you're short on time). Works whether running or paused.
- **Restart** — resets to seed data and starts again from beat 1.
- **Exit demo** — returns to the pre-start title card and clears the scripted incident. Auth and the real apps are untouched.
- Top-right links (Pilgrim / Volunteer / Management) open the real manual product in a new view if a judge asks to 'drive it yourself'.

If something looks off

- Nothing is advancing — check it isn't Paused (control bar says · Paused). Hit Resume, or Skip step.
- Panels look stale — hit Restart; it re-seeds from scratch.
- You reloaded mid-run — the timer chain can't survive a reload; hit Run Live Demo again (it resets first).
- Projector cut out — Pause, fix the display, Resume. The current beat is still on screen.

Timing

Full auto run is about 3 min 40 s. With a 30-second intro talk and a 20-second close you are just under 5 minutes. If you must trim live: after beat 12 (the 'why is this zone at risk?' breakdown), use **Skip step** through 13–15 while narrating the decision, then let 16–20 play.

5 · What is real, what is scripted

Genuinely real

Sarvam integration in the real apps: chat (Sarvam-M 105B) with native tool-calling, Mayura translation, Bulbul TTS, STT, and the existing fallback. The whole product architecture — ground reports, corroboration, Kumbh Pulse signal model, red-zone explainability, incident lifecycle, dispatch, task lifecycle, Event Log, offline behaviour, role permissions and the human-confirmation gates.

Scripted for the demo

The /demo timeline itself. Each of the 20 beats is a fixed call on the same store actions the real apps use, so the walkthrough always tells the same story in the same order. The AI text shown in /demo (assistant reply, translation line, Ops Copilot reasoning) is fixed copy, labelled 'scripted demo narration — deterministic, not a live generation'. The demo never calls Sarvam, so it can't stall, diverge, or fail on a network drop.

Never claimed

No real CCTV, live crowd counts, government feeds, emergency deployment, or real Kumbh incident data. Every operational number in /demo is synthetic and labelled 'Data: Synthetic / Simulation'. The pilgrim in the story is explicitly 'demo scenario — not a real pilgrim'.

Honesty lines to have ready

- “The demo is scripted so it's reliable on stage. The AI and the workflow underneath it are the real ones — open any of the three apps and drive it yourself.”
- “Every number here is synthetic. We are not claiming a live feed from the Kumbh.”
- “Sarvam is genuinely integrated — 105B chat with tool-calling, Mayura translate, Bulbul TTS. We just don't put a live model call on the critical path of a stage demo.”

6 · Judge Q&A prep

Is this AI making the decisions?

No. Every high-impact action — filing a report, promoting a signal, dispatching a volunteer — is behind a human confirmation. Setu and Ops Copilot propose; a person confirms. Kumbh Pulse is decision support and never triggers an automated action.

Why should we trust the risk score?

Because it shows its inputs. The 'why is this zone at risk?' panel breaks the score into four weighted factors — crowd activity, movement change, checkpoint reports, historical pattern — and lists the field signals behind them. It's explainable, not a black-box number.

Is the Sarvam integration real or a mock?

Real. The three apps use Sarvam-M 105B for chat with native tool-calling, Mayura for translation and Bulbul for TTS, with a graceful fallback. The /demo screen is deliberately scripted so a stage run can't fail on a bad network — that's a demo-reliability choice, not a capability gap.

What happens with no connectivity?

The pilgrim SOS path has an offline SMS fallback with cached first-response SOPs; ground reports queue offline and sync later. The /demo walkthrough runs entirely offline.

Where is the data stored? Is this multi-user?

For the prototype it's an in-memory store per browser, with optional cross-tab sync so the three roles stay in step when opened side by side. No database, no accounts — by design for the pitch. The data-access layer is a single module, so swapping in a real backend doesn't touch the screens.

How is a lost child handled specifically?

In-language guidance to stay put and go to the nearest Help Desk; a structured lost-person report; corroboration with other field reports and any found-child reports; a Critical incident with a language-matched responder and a translation bridge; and a full audit trail for reunification.

Could two reports be the same person double-counted?

That's the point of the corroboration step — the operator sees both source reports and decides. The signal raises it for a human; it doesn't auto-merge or auto-escalate.

What's the difference between Setu AI and Ops Copilot?

Same reasoning layer, different surface and role. Setu AI is for pilgrims and volunteers on the ground (guidance, translation, structuring a report). Ops Copilot is for the control room (recommend a responder, reason over the operational picture). Both propose; humans confirm.

7 · Facts, figures and IDs used in the demo

Scenario	Lost child — 6-year-old boy, red shirt, last seen near the Ramkund steps, Ghat 4 (Trimbakeshwar temple perimeter).
Pilgrim language	Tamil. Message: “என் குழந்தையை காணவில்லை” = “My child is missing.”
Ground reports	GR-4001 (V-218, M. Joshi, severity High) and GR-4002 (V-241, R. Kamble, severity Moderate).
Emerging signal	‘Multiple lost-person reports — Ghat 4’, confidence 69%, Kumbh Pulse decision support (synthetic).
Zone risk	Ghat 4 — score 84 / 100, band RED, forecast Critical in 10–20 min. Factors: Crowd Activity 42%, Movement Change 26%, Checkpoint Reports 18%, Historical Pattern 14%. Model pulse-rule-v0.3.
Incident	KS-1003 · Lost Person · Critical · Ghat 4 (promoted from GR-4001, then triaged).
Responder	V-233 — K. Bhosale, available, crowd-marshall trained, ~120 m from Ramkund, ETA ~4 min; Tamil translation bridged for the family.
Event Log chain	report created ×2 -> corroborated -> promoted to incident -> incident reported -> triaged -> volunteer dispatched -> task accepted -> arrived -> resolved.

Technical notes for the team

- Branch **feat/live-demo-mode**, pushed to origin. Two commits. **Not merged to main**.
- New files: `src/lib/demoScript.ts` (all scripted copy + the 20-beat table), `src/components/demo/DemoShared.tsx`, `src/components/demo/DemoStage.tsx`. Rewired: `DemoApp` / `DemoControls` / the three demo panes.
- Engine: a deterministic step-runner in the store (`runDemoStep` / `advanceDemo`) that calls existing actions. New store actions: `pauseDemo`, `resumeDemo`, `skipDemoStep`, `restartDemo`, `exitDemo`.
- One core change, backward-compatible: `dispatchIncident` takes an optional `forceVolunteerId` so the control room (and the demo) can assign a named responder. Default behaviour unchanged.
- Gates green: tsc clean; eslint clean (no new warnings); `npm run eval:setu 41/41`; next build passes.
- Verified in-browser: all 20 beats; Pause / Resume / Skip / Restart / Exit; unattended auto-advance; no console or hydration errors; Sarvam status still all-true; login / auth flow untouched.

Kumbh Setu — Live Demo master document. Generated for the team. All operational values are synthetic.